



Optimizing application performance through optimizing compilation

Francois Flandin

Encadré par
Pr Sid Touati

First semester of year 2024-2025

Objective

Explore how compilers optimizes programs using several optimizations levels : -O0 , -O1 , -O2 , -O3 , -Os .

2 parts to the project

1. Compiling 2 programs in C /C++ with each optimization level and compiler to compare performances
2. Checking which optimization is enabled for each optimization level and compiler

Introduction

Experience

- Environnement

- Method

Results

- Matrix Multiplication

- Dijkstra

Compilers Optimization

- How to get Compilers' Optimizations?

- Common optimizations on gcc , clang , icx

- What about ccomp?

Conclusion

Introduction

Experience

Environnement

Method

Results

Matrix Multiplication

Dijkstra

Compilers Optimization

How to get Compilers' Optimizations?

Common optimizations on gcc , clang , icx

What about ccomp?

Conclusion

Configuration

Computer in a lightweight configuration, avoid OS's optimizations and bloat from other programs or graphical interface.

OS and compilers

OS : Fedora Linux Workstation v40

gcc : version 14.2.1

icx : version 2024.2.1

clang : version 18.1.8

ccomp : version 3.14

Programs

There are 2 programs to compile : Matrix Multiplication (C) and Dijkstra's algorithm (C++). For each compiler and optimization level.

Object Size

The matrix size is set to two times the size of the largest cache.
The graph size is really huge to have a significant execution time.

Measurement

Function `gettimeofday()` placed before and after the main computation, measuring initialisation time is not the goal.

Finally

Run each program 12 times to visualize the data with R

Introduction

Experience

Environnement

Method

Results

Matrix Multiplication

Dijkstra

Compilers Optimization

How to get Compilers' Optimizations?

Common optimizations on gcc , clang , icx

What about ccomp?

Conclusion




Figure – Evolution of the execution time of the program `mat_mult.c` as a function of compiler and optimization level.




Figure – Evolution of the execution time of the program `dijkstra.c` as a function of compiler and optimization level.

	gcc	clang	icx
total time taken	136.37s	146.98s	155.65s
dijkstra usage	$\approx 19\%$	$\approx 9\%$	$\approx 11\%$
init/free usage	$\approx 81\%$	$\approx 91\%$	$\approx 89\%$

	gcc	clang	icx
total time taken	136.37s	146.98s	155.65s
dijkstra usage	$\approx 19\%$	$\approx 9\%$	$\approx 11\%$
init/free usage	$\approx 81\%$	$\approx 91\%$	$\approx 89\%$

C++ type inference

```
for (auto edge : temp_neighbors)
```

Introduction

Experience

- Environnement

- Method

Results

- Matrix Multiplication

- Dijkstra

Compilers Optimization

- How to get Compilers' Optimizations?

- Common optimizations on gcc , clang , icx

- What about ccomp?

Conclusion

gcc

```
gcc --help=optimizers -Q -On
```

clang

```
clang -On -emit-llvm -S program.c -o program.ll  
opt -On -debug-pass=manager program.ll -o program.ll
```

icx and ccomp

Everything is inside their documentation.

Peephole optimizations

Optimize a program over a reduced windows that slides through the program.

Peephole optimizations

Optimize a program over a reduced windows that slides through the program.

Strength reduction

Replace heavy computations with lighter ones.

Peephole optimizations

Optimize a program over a reduced windows that slides through the program.

Strength reduction

Replace heavy computations with lighter ones.

Code motion

Allow the compiler to move instructions around the program.

Peephole optimizations

Optimize a program over a reduced windows that slides through the program.

Strength reduction

Replace heavy computations with lighter ones.

Code motion

Allow the compiler to move instructions around the program.

Elimination of common subexpressions

Reuse previously computed values.

Dead argument elimination

Removes unused function arguments.

Dead argument elimination

Removes unused function arguments.

Inlining

Replace function calls with function's body.

Dead argument elimination

Removes unused function arguments.

Inlining

Replace function calls with function's body.

Constant propagation

Replace constant variable calls with the value.

Dead argument elimination

Removes unused function arguments.

Inlining

Replace function calls with function's body.

Constant propagation

Replace constant variable calls with the value.

Common subexpression elimination

Reuse previously computed values.

Loop unrolling

Reduces loop overhead by duplicating loop's body several times.

Loop unrolling

Reduces loop overhead by duplicating loop's body several times.

Loop interchange

Improve memory access by modifying the loops order in a serie of nested loops.

Loop unrolling

Reduces loop overhead by duplicating loop's body several times.

Loop interchange

Improve memory access by modifying the loops order in a serie of nested loops.

Loop peeling

Simplifies loops by removing non changing computation out of the loop.

Loop unrolling

Reduces loop overhead by duplicating loop's body several times.

Loop interchange

Improve memory access by modifying the loops order in a serie of nested loops.

Loop peeling

Simplifies loops by removing non changing computation out of the loop.

Loop jamming/fusion

Combine two adjacent loops into a single loop.

Does not activate much, but it has its optimizations!

Constant propagation

Replace constant variable calls with the value.

Does not activate much, but it has its optimizations !

Constant propagation

Replace constant variable calls with the value.

Common subexpression elimination

Reuse previously computed values.

Does not activate much, but it has its optimizations !

Constant propagation

Replace constant variable calls with the value.

Common subexpression elimination

Reuse previously computed values.

Inlining

Replace function calls with function's body.

ccomp

ccomp focuses on formally verified code translation so optimizations and performance is not a question.

gcc

gcc is the most efficient in general-purpose tasks and nested loops in C .

clang

clang was the most performant on C++ programs using specific features like type inferences.

icx

icx not being the most performant on this configuration could mean that the compiler is made for other types of workloads.

Each compiler has its own purpose, so it's important to choose the right compiler for a specific task.

What next ?

Further exploration on *energy consumption* optimization, in line with the model of Apple, a balance between *power* and *efficiency*.