

TER REPORT

---

# Optimizing application performance through optimizing compilation

---

*Prepared by*  
**Francois Flandin**

*Supervised by*  
Pr Sid Touati



# Contents

<b>1</b>	<b>Introduction</b>	<b>2</b>
<b>I</b>	<b>The experience</b>	<b>3</b>
<b>2</b>	<b>Method of experiment</b>	<b>4</b>
<b>3</b>	<b>Results</b>	<b>4</b>
3.1	Matrix Multiplication Results . . . . .	4
3.2	Dijkstra Results . . . . .	5
<b>II</b>	<b>Compilers</b>	<b>7</b>
<b>4</b>	<b>gcc</b>	<b>7</b>
4.1	Optimization options . . . . .	7
4.1.1	Method of obtention . . . . .	7
4.1.2	Options introduced by -O0 . . . . .	7
4.1.3	Options introduced by -O1 . . . . .	7
4.1.4	Options introduced by -O2 . . . . .	8
4.1.5	Options introduced by -O3 . . . . .	8
4.1.6	Options enabled by -Os . . . . .	9
<b>5</b>	<b>icx</b>	<b>9</b>
5.1	Optimization options . . . . .	9
5.1.1	Method of obtention . . . . .	9
5.1.2	Options introduced by -O0 . . . . .	9
5.1.3	Options introduced by -O1 . . . . .	9
5.1.4	Options introduced by -O2 . . . . .	9
5.1.5	Options introduced by -O3 . . . . .	10
<b>6</b>	<b>ccomp</b>	<b>10</b>
6.1	Optimization options . . . . .	10
6.1.1	Method of obtention . . . . .	10
6.1.2	Options introduced by -O0 . . . . .	10
6.1.3	Options introduced by -O1 , -O2 , -O3 . . . . .	10
6.1.4	Options introduced by -Os . . . . .	11
<b>7</b>	<b>clang</b>	<b>11</b>
7.1	Optimization options . . . . .	11
7.1.1	Method of obtention . . . . .	11
7.1.2	Options enabled at -O0 . . . . .	12
7.1.3	Options introduced by -O1 . . . . .	12
7.1.4	Options introduced by -O2 , -O3 , -Os . . . . .	13
<b>8</b>	<b>Conclusion</b>	<b>13</b>

# 1 Introduction

The tutoring project titled "Optimizing Application Performance through Compilation Optimization" explores how compilers improve the generated code using various optimization levels: `-O0` , `-O1` , `-O2` , `-O3` , and `-Os` .

The different optimization levels goes from `-O0` , which means no optimization, to `-O3` , which enables the most optimizations, but also `-Os` , which compiles to have the most compact code.

This project is divided into two main phases:

- Experimentation: Compile two programs with four different compilers (`gcc` , `icx` , `clang` , and `ccomp` ) at each optimization level, measure execution times over 12 iterations, and plot the data using `R`.
- Optimization Analysis: Study the optimization options enabled by each compiler at each level, to better understand their internal mechanisms.

These steps allow for the evaluation of compilers not only in terms of speed but also by identifying the specific characteristics of each optimization level.

## Part I

# The experience

## Experience's environnement

In this part will be presented the environnement of the experimentations, on which hardware it was done, on which software, etc.

### Hardware

In this part, the hardware part of the computer will be explored.

To explore the hardware of a computer on linux, the folder `/cpu` can be explored, it gives us the *topology* of the cpu, like which threads are on each cores, which caches are on each cores, discover what is the cache line size etc...

But there are also a lot of linux tools that allows to check the hardware's specifications of a computer, and basically doing the job for us.

There are graphical tools such as `likwid`, which provides the command `likwid-topology`, or the tool `hwloc`, which provides `lstopo`, a command that gives a beautiful GUI of the computer's hardware's specification.

**Model name :** 11th Gen Intel(R) Core(TM) i5-1135G7 @ 2.40GHz

**Adress size :** 39 bits physical, 48 bits virtual

**Cache line size :** 64 bytes

**Cores :** 4

Graphical Topology				
Coeurs	0 4	1 5	2 6	3 7
Cache L1	48 kB	48 kB	48 kB	48 kB
Cache L2	1MB	1MB	1MB	1MB
Cache L3	8 MB			

Table 1: Computer's topology

### Software

For the software part, the experiments were done on a lightweight configuration of the computer, where all unnecessary services were disabled, including the graphical interface, this allows the computer to run as baremetal as possible.

Here is a description of all the software elements used in this benchmark :

**OS** Fedora Linux v40 WorkStation

**gcc** version 14.2.1 20240912

**icx** version 2024.2.1.20240711

**clang** version 18.1.8

**ccomp** version 3.14

## 2 Method of experiment

To compare the performance of the different compilers, two programs will be used, the first `matrix_multiply`, is a C program that multiplies two matrices to put the result in a third matrix, it is an interesting challenge for optimization because of the presence of 3 nested loops for the calculation. This program will be compiled using `gcc`, `clang`, `icx` and `ccomp`.

The second program is `dijkstra` algorithm using C++, the goal behind this is to compare the compilers optimizations when the program has a lot of memory usage, because it is in C++, the program will only be compiled with the C++ versions of `gcc`, `clang` and `icx`, due to the lack of C++ compilation from `ccomp`.

Next, we can measure time taken by the program by placing the c-function `gettimeofday()` around the function that we need to calculate, initialisation measurement is not the main goal, but it can be measured too. The programs are executed 12 times for each optimization level to make an interesting violin plot using R.

## 3 Results

### 3.1 Matrix Multiplication Results




Figure 1: Evolution of the execution time of the program `mat_mult.c` as a function of compiler and optimization level.

The results for the C program (matrix multiplication) highlight clear differences in how each compiler handles optimizations.

`gcc` stands out as the best-performing compiler for this benchmark. Starting slower at `-O0`, it

quickly improves with each optimization level, reaching the best execution times at `-O2` and `-O3`. This suggests that `gcc` applies transformations that efficiently address the nested loop structure of matrix multiplication, making it highly suitable for such tasks.

`clang` shows a reliable improvement across optimization levels, but it remains behind `gcc`. While its performance gains are noticeable, it doesn't manage to handle the program as efficiently, especially at higher optimization levels. This difference could point to a less aggressive approach to handling the bottlenecks present in nested loops.

Despite being developed for Intel architectures, `icx` does not perform as well as expected. Its optimizations do improve execution times, but not enough to compete with `gcc` or `clang`. This makes `icx` less effective for matrix-heavy computations, suggesting it might prioritize other types of workloads.

`ccomp` delivers the slowest execution times in this benchmark, which is in line with its goal of producing formally verified code. While correctness is ensured, performance is not a priority, and its constant execution times across optimization levels reflect this limitation.

## 3.2 Dijkstra Results




Figure 2: Evolution of the execution time of the program `dijkstra.cpp` as a function of compiler and optimization level.

What we can notice from these 2 graphs, is how differently `gcc` compares with `clang` and `icx` according to the language used, C or C++. When using C, `gcc` produces way more optimized codes than `clang` and `icx`, whereas in C++, `gcc` is the least performant one.

Few experiments were done to get some answers to why it is like this.

First, the experiments were made on the C++ programs, at the compilation level of `-O3`, this will give us some data to work with, and, by using the command `perf record` and `perf report`, we can use material performance counters to get a lot of data on the percentage of time used by a function, we could use profiling applications such as `gprof`, but `perf` allowed us to have more precise data and avoided us to recompile entirely the program.

The results were the following:

	gcc	clang	icx
total time taken	136.37s	146.98s	155.65s
dijkstra usage	$\approx 26s$	$\approx 13s$	$\approx 17s$
init/free usage	$\approx 109s$	$\approx 133s$	$\approx 138s$

What we can get as results from here is that `gcc` is the most optimized in terms of initialisation and memory management, but is the worst one in terms of pure dijkstra calculation. If we push further the analysis, we can see there are some points in the program where types inference were used, such as `for (auto edge : temp.neighbors)`. `gcc` might not be able to manage this type system as well as `clang` or `icx`, hence the lack of optimizations.



## Part II

# Compilers

In this part will be presented all the compilers used in this tutorial project, as well as deepening their optimization options for each optimization levels.

The optimizations presented are not exhaustive, only the most important optimizations will be presented, and the method to obtain the optimizations will be given for each compiler, so the information can be obtained.

## 4 gcc

gcc is the most popular C compiler, it comes in the GNU Compiler Collection, which includes compilers for Fortran, C++, Go...

### 4.1 Optimization options

#### 4.1.1 Method of obtention

gcc is really helpful regarding getting information about what he does, because the compiler has an option for this sole purpose, which is

```
gcc --help=optimizers -Q -On
```

#### 4.1.2 Options introduced by -O0

gcc -O0 was originally meant to disable every optimization flag to give the least optimized code, however, the competition with other compilers led to now have some optimization flags enabled, 53 to be more precise. Not all optimizations will be presented, because the "optimizations" are for the most, analysis options.

A part of the optimizations concerns loop optimizations, with the option `-faggressive-loop-optimizations`, gcc enables, as its name suggests, a lot of loop optimizations, such as **loop fusion** and **fission**, **loop unrolling**, **loop peeling**.

Next is the dead code elimination, gcc can also remove every line of code that isn't used, this doesn't impact the execution speed, but reduces the size of the generated code.

And finally are the peephole optimizations this option introduces optimizations on a reduced window, optimizing redundant register moves, redundant calculations, etc.

Other enabled options at -O0 are either analysis options or add error-handling instructions.

#### 4.1.3 Options introduced by -O1

##### Intermediate Tree Optimizations

gcc -O1 optimizations add optimizations to the tree that the compiler builds in the intermediate level.

First, there is the dead code and redundancy elimination, with the options `-ftree-dce`, `-ftree-fre` and `-ftree-dse`, gcc removes all the **dead code**, computations that produce the same results as well as all **dead stores**, that means that the cpu has less operations to manage, and the code size is reduced.

Secondly, there is value and copy propagation, first, the option `-ftree-ccp` tells the compiler to propagate constant values across conditions when it's possible, it simplifies all the control flow

and makes place for further optimizations, and can even help for **branch prediction**.

Secondly, the option **-ftree-copy-prop** propagates values with copy instructions, for example `b = a; c = b;` would be replaced with `c = a;`.

And finally, loop optimizations, there are two main options here to consider, the first, **-ftree-sink** moves computations outside of the loop when it's possible, it allows to avoid computing the same thing over and over.

The second is **-ftree-sra** for *scalar replacement of aggregates*, it breaks structures and array into individual variable when needed, it can improve register usage and enable even further optimizations. An example of this would be a variable of the struct `Point {int x; int y}` would be transformed from `Point p;` to `int p.x, p.y;`.

At **-O1**, there are also other options to enable inlining of functions, that means that a function call will be replaced by the function's body, which allows for further optimizations.

#### 4.1.4 Options introduced by -O2

##### Interprocedural Analysis Optimizations

gcc **-O2** introduces IPA (Interprocedural Analysis) optimizations, which aim to improve performance by analyzing and optimizing across function boundaries. This includes **-fipa-icf** (Identical Code Folding), which detects and merges functions or read-only variables with identical implementations, reducing code size and memory usage.

Additionally, **-fipa-cp** (Constant Propagation) propagates constant values across function calls, allowing the compiler to simplify computations and eliminate unnecessary branches. These optimizations are particularly effective in programs with repetitive patterns or predictable input values, contributing to leaner and faster executables.

##### Aligning

gcc **-O2** introduces aligning of functions, jumps, labels, and loops, the aligning allows to ensure that the instructions fit at the beginning of cache lines, which can reduce cache misses. Also, aligning labels and jumps allows to decrease branch prediction misses.

**-O2** also enables **-funroll-loops**, which performs loop unrolling when iteration count is known, which allows the compiler to optimize as it wishes.

#### 4.1.5 Options introduced by -O3

##### Loop Optimizations

gcc **-O3** enables mostly further loops optimizations, such as **loop peeling**, which removes special cases in the loop and put them outside of the loop, then, **loop jamming** combines two loops with the same scope/range into a single one, enabling more optimizations, and is generally combined with **loop unrolling**.

**loop unrolling** reduces loop overhead by duplicating the loop body multiple times, minimizing the number of iterations and improving execution efficiency, also **loop interchange**, swap two loops in a nested loop, this allows to put the most changing variable to the consecutive elements in memory, instead of having the most changing variable to separate elements in memory, it's very useful in programs such as **matrix multiplication**.

Finally, **loop splitting** allows simplification of a loop by dividing it into several loops, it simplifies the loops with the pattern `for (int i = 0; i < 10; i++) if(i < 5)... else...` into two loops, one with the range 0..4 and the other 5..9. It helps in reducing dependencies.

#### 4.1.6 Options enabled by `-Os`

`gcc -Os` enables the same options as `gcc -O2` , but removes the aligning optimizations to free some space.

## 5 `icx`

It's Intel's compiler for their x86 architecture, it is made to deliver extremely optimized programs on their processor architecture.

### 5.1 Optimization options

#### 5.1.1 Method of obtention

Everything is available in the online documentation, it's relatively easy to find what we need.

#### 5.1.2 Options introduced by `-O0`

Everything is disabled at `-O0` .

#### 5.1.3 Options introduced by `-O1`

At this level of optimization, `icx` starts by introducing **data flow analysis**, to help gathering data about data flows throughout the whole program, allowing for optimizations.

Then, the compiler can also reorganize code as he wishes with **code motion**, which allows to move non-changing result operations out of loop to avoid recomputing it over and over, and **instruction scheduling**, which allows to change the order of instructions, to hide memory latency for exemple, or doing something while something heavier is done. It is done to reduce CPU stall.

Other optimizations are also enabled, **strength reduction**, that changes heavy operations for cheaper ones, like switching from a multiplication to a bit shift, **test replacement**, that simplifies loops' conditions for exemple :

```
if (x >= 0 && x < 100) { /* code */ } // base loop
if ((unsigned)x < 100) { /* code */ } // optimized loop
```

And, finally, **split-lifetime analysis**, can split the *lifetime* of a variable into smaller lifetimes, to reallocate the registers when the variable is not used anymore.

#### 5.1.4 Options introduced by `-O2`

At this level of optimization, `icx` introduces more optimizations, like **constant propagation**, or **forward substitution**, which can replace variable with their expressions, for exemple :

```
int a = 5 + i
int b = a + 3
```

will be replaced by

```
int b = 5 + i + 3
```

These two options complete each other really well, because imagine that the variable `i` was a constant, it would be replaced by its value, hence further optimizing the program.

Additional optimizations are introduced, like **loop unrolling**, which reduces the number of iterations by executing multiple iterations within a single loop cycle, **peephole optimizations**, which examines small code fragments to identify and replace inefficiencies, such as redundant instructions or suboptimal sequences, improving code quality. Or even **optimized code selection**, replaces inefficient instructions or code sequences with more efficient alternatives.

There are also code-removing optimizations such as **dead-code elimination**, **dead-store elimination**, and **dead static function elimination** which does not really improve performance but can help reduce the weight put on the CPU.

### 5.1.5 Options introduced by `-O3`

It enables all the same optimizations as `-O2` but add more aggressive loop transformations such as **loop fusion**, **block-unroll-and-jam**, with these two options, the compiler expands and finds the best fusions/jams in the loops and **collapsing-IF-statements**, where multiple successive ifs are combined into one to reduce branching therefore gaining performances.

## 6 *ccomp*

### *Certified Compiler*

It's the only C compiler that is formally verified to produce a code described by the source code, optimized or not. By its formally verified nature, it is the one that produces the most efficient code. It is generally used for safety critical programs where a bug introduced by compilation could lead to serious problems.

## 6.1 Optimization options

### 6.1.1 Method of obtention

Everything is available in the online documentation, it's relatively easy to find what we need.

### 6.1.2 Options introduced by `-O0`

`ccomp` doesn't enable anything at `-O0`, and the main optimization level `-O` gives the same optimizations as `-O1`, `-O2`, `-O3`, because they are just an alias of that optimization option.

### 6.1.3 Options introduced by `-O1`, `-O2`, `-O3`

There are not a lot of options so here is a list of them

**-fcont-prop** enables constant propagation, a process that replaces variables with their known constant values throughout the program.

**-fcse** activates the elimination of common subexpressions, reducing redundant calculations by reusing previously computed values.

**-fif-conversion** with this option, the compiler decides whether to replace simple if-then-else statements or the conditional operator (?:) with conditional assignments.

The heuristic-based approach selectively optimizes small and balanced expressions, provided the target architecture supports a suitable conditional move instruction.

**-finline** enables or disables the inlining of functions, potentially improving performance by replacing function calls with the actual function code.

**-finline-functions-called-once** specifically inlines functions that are called only once, reducing overhead while maintaining efficiency.

**-fredundancy** activates the elimination of redundant computations and unnecessary memory stores, improving execution time by avoiding repetitive operations.

**-ftailcalls** optimizes function calls in tail position, which can enhance performance by reusing the current function's stack frame.

**-ffloat-const-prop 2** enables full propagation of floating-point constants, ensuring arithmetic is performed in IEEE double precision format with round-to-nearest mode.

#### 6.1.4 Options introduced by -Os

The documentation doesn't give a lot of informations about it. Just the usual "optimizes for code size".

## 7 clang

It's a compiler made by the LLVM Developer Group, it can replace `gcc` by supporting most of the option flags of `gcc`.

The objective behind this compiler is to create a compiler that allowed: first, better diagnostics, which allows to debug more easily, second, to be separated from the GNU licence, which forced softwares to be integrated to said licence and therefore having to open-source proprietary software, third, to have a nimble compiler that is simple and easy to develop and maintain.

### 7.1 Optimization options

#### 7.1.1 Method of obtention

`clang`'s documentation does not explicitly specify all enabled optimizations, however, it has some commands that can give what is enabled with each optimization level, there are 2 commands because of the way `clang` works.

The compiler is divided in two parts, `clang`, the part where the intermediate representation (IR) is made, and `opt`, the part that optimizes said IR. The first command is used to know which optimizations are enabled by `opt`.

```
llvm-as < /dev/null | opt -On -disable-output -debug-pass-manager
```

The second command gives the optimizations added by `clang` itself

```
diff -wy --suppress-common-lines \  
  <(echo 'int;' | clang -xc - -o /dev/null -\#\#\# 2>&1 | tr " " "\n" | grep -v /tmp) \  
  <(echo 'int;' | clang -xc -On - -o /dev/null -\#\#\# 2>&1 | tr " " "\n" | grep -v /tmp)
```

Another way of doing this, which gives way more optimizations is the combination of the commands

```
clang -O2 -emit-llvm -S matrix_multiply.c -o mat_mult.ll
opt -O2 -debug-pass=Manager mat_mult.ll -o mat_mult.ll
```

### 7.1.2 Options enabled at -O0

clang like the other compilers doesn't enable a lot of optimizations, the optimizations are called **AlwaysInlinerPass**, which inlines functions marked with the `always_inline` attribute, **CoroConditionalWrapper**, which is related to coroutines support in C++, which is a function that can be paused and resumed, enabling asynchronous and cooperative task management. Finally, **VerifierPass** ensures that the LLVM Intermediate Representation generated by the compiler is correct and coherent with the original program.

### 7.1.3 Options introduced by -O1

#### Loop Optimizations

At -O1, clang introduces several important loop optimizations to enhance runtime performance. The **LoopSimplifyPass** ensures that loops are structured in a canonical form, making them easier to analyze and optimize in subsequent passes. **LCSSAPass** (Loop-Closed SSA) further prepares loops by transforming them into a representation that facilitates optimizations involving memory access and data dependencies. **IndVarSimplifyPass** refines induction variables, such as loop counters, simplifying loop conditions and arithmetic to improve execution flow. Meanwhile, **LoopRotatePass** adjusts loop headers, improving branch prediction and enabling opportunities for more aggressive transformations like unrolling in later optimization stages. Though full unrolling and vectorization are not fully activated at this level, these preparatory transformations lay the groundwork for higher-level optimizations while maintaining reasonable compilation times.

#### Coroutines

For programs utilizing coroutines, clang introduces key transformations to handle their unique execution model. **CoroEarlyPass** prepares coroutine constructs by splitting their code into manageable parts for further optimization. **CoroElidePass** eliminates unnecessary coroutine frames when they are not required, reducing runtime overhead. Finally, the **CoroCleanupPass** finalizes coroutine transformations by removing temporary constructs and generating efficient, optimized code for coroutine-based workflows. These coroutine-specific passes ensure that clang can handle modern programming paradigms efficiently, even at -O1.

#### Global Optimizations

clang's -O1 also includes broader global optimizations that significantly impact performance and code quality. **IPSCCPPass** (Interprocedural Sparse Conditional Constant Propagation) enables the propagation of constant values across function boundaries, revealing opportunities for eliminating redundant computations and simplifying control flow. The **LICMPass** (Loop-Invariant Code Motion) complements loop optimizations by hoisting computations that remain constant across iterations, reducing repeated calculations. Additionally, **DeadArgumentEliminationPass** simplifies function signatures by removing unused arguments, and **MemCpyOptPass** optimizes memory operations like `memcpy` by replacing them with more efficient alternatives where possible. Together, these optimizations strike a balance between improving runtime performance

and maintaining fast compilation speeds, providing a strong baseline for further enhancements at higher optimization levels.

#### 7.1.4 Options introduced by `-O2` , `-O3` , `-Os`

At `-O2` , `-O3` , and `-Os` , `clang` focuses on impactful optimizations like **GVNPass** (Global Value Numbering), which eliminates redundant computations, and **DSEPass** (Dead Store Elimination), which removes unnecessary memory writes.

**SLPVectorizerPass** enhances performance by vectorizing independent operations for SIMD, while **JumpThreadingPass** simplifies control flow by bypassing unnecessary conditions. Together, these key passes significantly improve execution efficiency and streamline code generation for both performance-critical and size-constrained applications.

## 8 Conclusion

This project has given us a better understanding of how various compilers handle optimizations and how these optimizations impact performance.

The experiments revealed that `gcc` is the top performer when it comes to optimizing nested loops in C programs, consistently delivering superior results. On the other hand, `clang` showed its strength with C++ workloads, surpassing `gcc` in tasks like Dijkstra’s algorithm thanks to its deep understanding of modern C++ features. Despite being designed for Intel architectures, `icx` did not perform well in highly computational scenarios, suggesting that its optimizations may be geared towards other types of workloads. `ccomp` , on the other hand, focused on formally verified code, highlighting the balance between correctness and performance and making it a unique but specialized tool for safety-critical applications.

In conclusion, it is essential to choose the right compiler for specific tasks. While `gcc` is versatile and reliable for general-purpose and performance-critical applications, `clang` shines in C++ workloads due to its modern language features support. `icx` and `ccomp` , although not leading in raw performance, cater to specific use cases, showing that each compiler has its strength in different scenarios.

This analysis opens the door for future exploration, such as investigating how compilers can optimize code to be more energy-efficient. Because since the beginning of computer science the major concern was pure performance and less energy-efficiency, this could be where technological advancements truly shine, not just in creating more powerful computers, but in developing more efficient ones. A prime example of this approach is Apple’s ecosystem, which achieves impressive performance without consuming excessive energy. Future innovations could follow this model, focusing on both power and efficiency.

## References

- Stack overflow question about clang optimization levels, 2013. URL <https://stackoverflow.com/questions/15548023/clang-optimization-levels>.
- LLVM Group. Clang compiler documentation, 2024. URL <https://clang.llvm.org/>.
- Intel Corporation. Intel® oneapi dpc++/c++ compiler developer guide and reference, 2024. URL <https://www.intel.com/content/www/us/en/docs/dpcpp-cpp-compiler/developer-guide-reference/2024-2/o-001.html>.
- Xavier Leroy. The compcert c verified compiler, documentation and user’s manual. version 3.14, 2024. URL <https://compcert.org/man/manual003.html#sec26>.